



Using Red Hat Enterprise Linux 6.4 with NetApp ONTAP

ONTAP SAN Host

Amanda Stroman, Sean Daffy
June 25, 2020

This PDF was generated from https://docs.netapp.com/us-en/ontap-sanhost/hu_rhel_64.html on June 26, 2020.
Always check docs.netapp.com for the latest.

Table of Contents

- Using Red Hat Enterprise Linux 6.4 with NetApp ONTAP..... 1
 - Installing the Linux Unified Host Utilities 1
 - SAN Toolkit..... 1
 - SAN Booting..... 2
 - Multipathing..... 2
 - Recommended Settings..... 4
 - Known Problems and Limitations 6
 - Release Notes..... 8

Using Red Hat Enterprise Linux 6.4 with NetApp ONTAP

Installing the Linux Unified Host Utilities

The NetApp Linux Unified Host Utilities software package is available on the [NetApp Support Site](#) in a 32-bit and 64-bit .rpm file. If you do not know which file is right for your configuration, use the [NetApp Interoperability Matrix Tool](#) to verify which one you need.

Installing the Linux Unified Host Utilities is strongly recommended, but not mandatory. The utilities do not change any settings on your Linux host. The utilities improve management and assist NetApp customer support in gathering information about your configuration.

Before you begin

If you have a version of Linux Unified Host Utilities currently installed you should upgrade it or, you should remove it and use the following steps to install the latest version.

1. Download the 32-bit or 64-bit Linux Unified Host Utilities software package from the [NetApp Support Site](#) to your host.
2. Use the following command to install the software package:

```
rpm -ivh netapp_linux_unified_host_utilities-7-1.x86_64
```

SAN Toolkit

The toolkit is installed automatically when you install the NetApp Host Utilities package. This kit provides the **sanlun** utility, which helps you manage LUNs and HBAs. The **sanlun** command returns information about the LUNs mapped to your host, multipathing, and information necessary to create initiator groups.

Example

In the following example, the **sanlun lun show command** returns LUN information.

```
# sanlun lun show all
controller(7mode/E-Series)/          device    host
vserver(cDOT/FlashRay)  lun-pathname  filename  adapter  protocol  lun
-----
data_vserver            /vol/vol1/lun1  /dev/sdb  host16    FCP        120.0g  cDOT
data_vserver            /vol/vol1/lun1  /dev/sdc  host15    FCP        120.0g  cDOT
data_vserver            /vol/vol2/lun2  /dev/sdd  host16    FCP        120.0g  cDOT
data_vserver            /vol/vol2/lun2  /dev/sde  host15    FCP        120.0g  cDOT
```

SAN Booting

Before you begin

If you decide to use SAN booting, it must be supported by your configuration. You can use the [NetApp Interoperability Matrix Tool](#) to verify that your OS, HBA, HBA firmware and the HBA boot BIOS, and ONTAP version are supported.

1. Map the SAN boot LUN to the host.
2. Verify multiple paths are available.

Remember, multiple paths will only be available after the host OS is up and running on the paths.

3. Enable SAN booting in the server BIOS for the ports to which the SAN boot LUN is mapped.

For information on how to enable the HBA BIOS, see your vendor-specific documentation.

4. Reboot the host to verify the boot is successful.

Multipathing

For Red Hat Enterprise Linux (RHEL) 6.4 the `/etc/multipath.conf` file must exist, but you do not need to make specific changes to the file. RHEL 6.4 is compiled with all settings required to recognize and correctly manage ONTAP LUNs.

To Enable ALUA Handler, perform the following steps:

1. Create a backup of the `initrd`-image.
2. Append the following parameter value to the kernel for ALUA and non-ALUA to work:

`rdloaddriver=scsi_dh_alua`

Example

```
kernel /vmlinuz-2.6.32-358.6.1.el6.x86_64 ro root=/dev/mapper/ vg_ibmx355021082-  
lv_root rd_NO_LUKS rd_LVM_LV=vg_ibmx355021082/ lv_root LANG=en_US.UTF-8  
rd_LVM_LV=vg_ibmx355021082/lv_swap rd_NO_MD SYSFONT=latarcyrheb-sun16 crashkernel=auto  
KEYBOARDTYPE=pc KEYTABLE=us rd_NO_DM rhgb quiet rdloaddriver=scsi_dh_alua
```

3. Use the `mkinitrd` command to recreate the `initrd`-image.

RHEL 6x and later versions use either:

The command: `mkinitrd -f /boot/ initrd-"uname -r".img uname -r`

Or

The command: `dracut -f`

4. Reboot the host.
5. Verify the output of the `cat /proc/cmdline` command to ensure that the setting is complete.

You can use the `multipath -ll` command to verify the settings for your ONTAP LUNs.

The following sections provide sample multipath output for a LUN mapped to ASA and non-ASA personas.

All SAN Array Configuration

For All SAN Array (ASA) configuration there should be one group of paths with single priorities. All the paths are Active/Optimized, meaning they are serviced by the controller and I/O is sent on all the active paths.

Example

The following example displays the correct output for an ONTAP LUN with four Active/Optimized paths:

```
# multipath -ll
3600a0980383034466b2b4a3775474859 dm-3 NETAPP,LUN C-Mode
size=20G features='4 queue_if_no_path pg_init_retries 50 retain_attached_hw_handle'
hwhandler='1 alua' wp=rw
|+- policy='round-robin 0' prio=50 status=active
| |- 1:0:8:1 sdb 8:16 active ready running
| `-- 2:0:8:1 sdd 8:48 active ready running
`+- policy='round-robin 0' prio=10 status=enabled
  |- 1:0:9:1 sdc 8:32 active ready running
  `-- 2:0:9:1 sde 8:64 active ready running
```

Note

Do not use an excessive number of paths to a single LUN. No more than 4 paths should be required. More than 8 paths might cause path issues during storage failures.

Non-ASA Configuration

For non-ASA configuration there should be two groups of paths with different priorities. The paths with the higher priorities are Active/Optimized, meaning they are serviced by the controller where the aggregate is located. The paths with the lower priorities are active but are non-optimized because they are served from a different controller. The non-optimized paths are only used when no optimized paths are available.

Example

The following example displays the correct output for an ONTAP LUN with two Active/Optimized paths and two Active/non-Optimized paths:

```
# multipath -ll
3600a0980383034466b2b4a3775474859 dm-3 NETAPP,LUN C-Mode
size=20G features='4 queue_if_no_path pg_init_retries 50 retain_attached_hw_handle'
hwhandler='1 alua' wp=rw
|+- policy='round-robin 0' prio=50 status=active
| |- 1:0:8:1 sdb 8:16 active ready running
| `-- 2:0:8:1 sdd 8:48 active ready running
`+- policy='round-robin 0' prio=10 status=enabled
  |- 1:0:9:1 sdc 8:32 active ready running
  `-- 2:0:9:1 sde 8:64 active ready running
```

Note

Do not use an excessive number of paths to a single LUN. No more than 4 paths should be required. More than 8 paths might cause path issues during storage failures.

Recommended Settings

The RHEL 6.4 OS is compiled to recognize ONTAP LUNs and automatically set all configuration parameters correctly.

The multipath.conf file must exist for the multipath daemon to start, but you can create an empty, zero-byte file using the command:

```
touch /etc/multipath.conf
```

The first time you create this file, you might need to enable and start the multipath services.

```
[root@jfs0 ~]#systemctl enable multipathd
[root@jfs0 ~]# systemctl start multipathd
```

There is no requirement to add anything directly to multipath.conf, unless you have devices that you do not want to be managed by multipath or you have existing settings that override defaults.

You can add the following syntax to the multipath.conf file to exclude the unwanted devices.

Note

Replace the <DevId> with the WWID string of the device you want to exclude. Use the following command to determine the WWID:

```
blacklist {
    wwid <DevId>
    devnode "^(ram|raw|loop|fd|md|dm-|sr|scd|st)[0-9]*"
    devnode "^hd[a-z]"
    devnode "^cciss.*"
}
```

Example

In this example, sda is the local SCSI disk that we need to blacklist.

1. Run the following command to determine the WWID:

```
# /lib/udev/scsi_id -gud /dev/sda
360030057024d0730239134810c0cb833
```

2. Add this WWID to the blacklist stanza in the `/etc/multipath.conf`:

```
blacklist {
    wwid 360030057024d0730239134810c0cb833
    devnode "^(ram|raw|loop|fd|md|dm-|sr|scd|st)[0-9]*"
    devnode "^hd[a-z]"
    devnode "^cciss.*"
}
```

You should always check your `/etc/multipath.conf` file for legacy settings, especially in the defaults section, that may be overriding default settings.

The table below shows the critical `multipathd` parameters for ONTAP LUNs and the required values. If a host is connected to LUNs from other vendors and any of these parameters are overridden, they will need to be corrected by later stanzas in `multipath.conf` that apply specifically to ONTAP LUNs. If this is not done, the ONTAP LUNs may not work as expected. These defaults should only be overridden in consultation with NetApp and/or OS vendor and only when the impact is fully understood.

Parameter	Setting
detect_prio	yes
dev_loss_tmo	"infinity"
failback	immediate
fast_io_fail_tmo	5
features	"3 queue_if_no_path pg_init_retries 50"
flush_on_last_del	"yes"
hardware_handler	"0"
no_path_retry	queue
path_checker	"tur"
path_grouping_policy	"group_by_prio"
path_selector	"round-robin 0"
polling_interval	5

Parameter	Setting
prio	"ontap"
product	LUN.*
retain_attached_hw_handler	yes
rr_weight	"uniform"
user_friendly_names	no
vendor	NETAPP

Example

The following example shows how to correct an overridden default. In this case, the `multipath.conf` file defines values for `path_checker` and `no_path_retry` that are not compatible with ONTAP LUNs. If they cannot be removed because of other SAN arrays still attached to the host, these parameters can be corrected specifically for ONTAP LUNs with a device stanza.

```
defaults {
    path_checker      readsector0
    no_path_retry     fail
}

devices {
    device {
        vendor        "NETAPP  "
        product        "LUN.*"
        no_path_retry  queue
        path_checker    tur
    }
}
```

Known Problems and Limitations

NetApp Bug ID	Title	Description	Bugzilla ID
673009	Creating ext4 file system on LV striped across 15 or more discard-enabled, thinly provisioned multipath devices triggers "request botched" kernel errors	"Request botched" kernel errors have been seen when users attempt to create an ext4 file system on discard-enabled, thinly provisioned multipath devices. As a result, creating the ext4 file system might take longer to complete and occasional disruption occurs. This issue has occurred only when users attempt to create the ext4 file system on a LV striped across 15 or more discard-enabled multipath devices on systems running Red Hat Enterprise Linux 6.x and Data ONTAP 8.1.3 and later operating in 7-Mode. The issue happens because the kernel erroneously attempts to merge discard requests, which is not supported on Red Hat Enterprise Linux 6.x at this time. When this issue occurs, multiple instances of the following message are written to syslog (/var/log/messages): kernel: blk: request botched. As a result, file system creation might take longer time to complete than expected.	907844

Release Notes

ASM Mirroring

ASM mirroring might require changes to the Linux multipath settings to allow ASM to recognize a problem and switch over to an alternate fail group. Most ASM configurations on ONTAP use external redundancy, which means that data protection is provided by the external array and ASM does not mirror data. Some sites use ASM with normal redundancy to provide two-way mirroring, normally across different sites. See [Oracle Databases on ONTAP](#) for further information.

Copyright Information

Copyright © 2020 NetApp, Inc. All rights reserved. Printed in the U.S. No part of this document covered by copyright may be reproduced in any form or by any means-graphic, electronic, or mechanical, including photocopying, recording, taping, or storage in an electronic retrieval system-without prior written permission of the copyright owner.

Software derived from copyrighted NetApp material is subject to the following license and disclaimer:

THIS SOFTWARE IS PROVIDED BY NETAPP “AS IS” AND WITHOUT ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE, WHICH ARE HEREBY DISCLAIMED. IN NO EVENT SHALL NETAPP BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

NetApp reserves the right to change any products described herein at any time, and without notice. NetApp assumes no responsibility or liability arising from the use of products described herein, except as expressly agreed to in writing by NetApp. The use or purchase of this product does not convey a license under any patent rights, trademark rights, or any other intellectual property rights of NetApp.

The product described in this manual may be protected by one or more U.S. patents, foreign patents, or pending applications.

RESTRICTED RIGHTS LEGEND: Use, duplication, or disclosure by the government is subject to restrictions as set forth in subparagraph (c)(1)(ii) of the Rights in Technical Data and Computer Software clause at DFARS 252.277-7103 (October 1988) and FAR 52-227-19 (June 1987).

Trademark Information

NETAPP, the NETAPP logo, and the marks listed at <http://www.netapp.com/TM> are trademarks of NetApp, Inc. Other company and product names may be trademarks of their respective owners.